

Project Report

Runchen Hu rh2619

Code: (Run in Spark Shell)

```
// SVM
.....
import org.apache.spark.mllib.regression.LabeledPoint
.....
import org.apache.spark.mllib.classification.SVMWithSGD
.....
import org.apache.spark.mllib.evaluation.BinaryClassificationMetrics

// First I standardized the dataframe and drop the rows with null values.
.....
val spark =
.....
org.apache.spark.sql.SparkSession.builder.master("local").appName("Reader").getOrCreate;
.....
val df = spark.read.format("csv").option("header", "true").load("titanic-train.csv")
.....
val data = df.select($"Survived", $"Pclass", $"Sex", $"Age", $"Fare")
.....
val newdata = data.withColumn("Sex", when(col("Sex") === "male",
.....
"1").otherwise(col("Sex")))
.....
val finaldata = newdata.withColumn("Sex", when(col("Sex") === "female",
.....
"0").otherwise(col("Sex")))
.....
val data2 = finaldata.selectExpr("cast(Survived as double) Survived", "cast(Pclass as int)
.....
Pclass", "cast(Sex as int) Sex", "cast(Age as int) Age", "cast(Fare as float) Fare")
.....
val data3 = data2.na.drop

// Then I split the dataset into training set and testing set.
.....
val splits = data3.randomSplit(Array(0.8, 0.2), seed = 11L)
.....
val training = splits(0).cache()
.....
val test = splits(1)

// I converted dataframe into RDD and set feature columns and label column.
.....
val training_rows: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = training.rdd
.....
val labeled_training = training_rows.map( a =>
.....
LabeledPoint(a.getDouble(0), org.apache.spark.mllib.linalg.Vectors.dense(a.getInt(1), a.getInt(2),
.....
a.getInt(3), a.getFloat(4))))
```



```
scala> data3.show
```

Survived	Pclass	Sex	Age	Fare
0.0	3	1	22	7.25
1.0	1	0	38	71.2833
1.0	3	0	26	7.925
1.0	1	0	35	53.1
0.0	3	1	35	8.05
0.0	1	1	54	51.8625
0.0	3	1	2	21.075
1.0	3	0	27	11.1333
1.0	2	0	14	30.0708
1.0	3	0	4	16.7
1.0	1	0	58	26.55
0.0	3	1	20	8.05
0.0	3	1	39	31.275
0.0	3	0	14	7.8542
1.0	2	0	55	16.0
0.0	3	1	2	29.125
0.0	3	0	31	18.0
0.0	2	1	35	26.0
1.0	2	1	34	13.0
1.0	3	0	15	8.0292

only showing top 20 rows

```
scala> val model = SVMWithSGD.train(labeled_training, numIterations)
17/12/07 20:44:18 WARN SVMWithSGD: The input data is not directly cached, which may hurt performance if its parent RDDs are also uncached.
17/12/07 20:44:19 WARN SVMWithSGD: The input data was not directly cached, which may hurt performance if its parent RDDs are also uncached.
model: org.apache.spark.mllib.classification.SVMModel = org.apache.spark.mllib.classification.SVMModel: intercept = 0.0, numFeatures = 4, numClasses = 2,
threshold = 0.0

scala> val metrics = new BinaryClassificationMetrics(scoreAndLabels)
metrics: org.apache.spark.mllib.evaluation.BinaryClassificationMetrics = org.apache.spark.mllib.evaluation.BinaryClassificationMetrics@9e5061e

scala> val auROC = metrics.areaUnderROC()
auROC: Double = 0.6033379058070417

scala> println("Area under ROC = " + auROC)
Area under ROC = 0.6033379058070417
```

Code: (Run in R)

```
library(e1071)
library(ROCR)

setwd('~Downloads/SVM') # You can set your own work directory

# First I standardized the original dataframe.
df = read.table('titanic-train.csv', header = TRUE, sep = ',')
head(df)
df2 = df[, -which(names(df) %in%
c("PassengerId", "Name", "SibSp", "Ticket", "Cabin", "Embarked", "Parch"))]
head(df2)
df3 = df2[complete.cases(df2),]
head(df3)

#Then I split df3 into two sets
set.seed(43)
tst_idx = sample(714, 200, replace = FALSE)
test = df3[tst_idx,]
training = df3[-tst_idx,]

X_training = training[, -which(names(training) %in% c("Survived"))]
y_training = training[,c("Survived")]
X_test = test[, -which(names(test) %in% c("Survived"))]
y_test = test[,c("Survived")]

#Then I did the AUC test.
aucs = c()
plot(x = NA, y = NA, xlim = c(0, 1), ylim = c(0, 1), ylab = 'True Positive Rate', xlab = 'False
Positive Rate', bty = 'n')

lvls = c(0, 1)
for (type.id in 1:2) {
  type = as.factor(training$Survived == lvls[type.id])
```

```
model = model <- svm(as.formula(paste("type", "~ .", sep = " ")), data = training[,-1], kernel
= "linear", probability = T, type = "C-classification",cost = 1)
```

```
Survived = predict(model, test[,-1], decision.values = T, probability = T)
score = attr(Survived, "probabilities")[,2]
```

```
actual.Survived = test$Survived == lvls[type.id]
pred = prediction(score, actual.Survived)
perf = performance(pred, "tpr", "fpr")
roc.x = unlist(perf@x.values)
roc.y = unlist(perf@y.values)
lines(roc.y ~ roc.x, col = type.id + 1, lwd = 2)
nauc = performance(pred, "auc")
nauc = unlist(slot(nauc, "y.values"))
aucs[type.id] = nauc
}
legend("bottomright", legend = c("1 vs others", "2 vs others"), col = 2:3, lty = 1, lwd = 2)
lines(x = c(0, 1), c(0, 1))
cat("AUC", fill = T)
aucs
```

ScreenShots:


```
> head(df3)
```

	Survived	Pclass	Sex	Age	Fare
1	0	3	male	22	7.2500
2	1	1	female	38	71.2833
3	1	3	female	26	7.9250
4	1	1	female	35	53.1000
5	0	3	male	35	8.0500
7	0	1	male	54	51.8625

```
> X_training
  Pclass  Sex  Age  Fare
1      3  male 22.00  7.2500
2      1 female 38.00 71.2833
4      1 female 35.00 53.1000
5      3  male 35.00  8.0500
7      1  male 54.00 51.8625
9      3 female 27.00 11.1333
11     3 female  4.00 16.7000
12     1 female 58.00 26.5500
15     3 female 14.00  7.8542
16     2 female 55.00 16.0000
```

```
> y_training
 [1] 0 1 1 0 0 1 1 1 0 1 0 0 1 1 0 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 0 0 1 0 1 0 0 1 0 0 0 0 0 1 0
[57] 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 1 1 1
[113] 1 1 1 0 0 1 1 0 1 0 1 0 0 1 0 1 0 0 1 0 1 0 0 0 0 0 0 1 0 0 0 1 1 1 1 0 0 0 0 1 1 1 1 0 0 0 0 1 0 1 0
[169] 1 0 1 0 1 1 0 1 1 0 1 0 0 1 1 1 1 1 0 0 1 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 1 1
[225] 0 1 1 0 0 0 0 1 0 1 1 0 0 0 0 0 0 0 1 1 1 0 1 0 0 1 0 0 1 1 1 1 0 1 1 0 1 0 0 0 0 0 1 1 0 0 0 0 1 0 1 1
[281] 0 0 0 1 1 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 1 1 1 0 0 1 0 1 0 1 0 1 0 0 0 0 1 1 0 1 1 0 0 1 0 1 0 0 1
[337] 0 0 1 0 1 0 0 1 0 0 0 0 1 0 1 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 1 0 0 1 0 1 0 0 0 1 1 1 0 0
[393] 0 1 1 1 1 1 0 0 0 1 0 1 0 0 1 0 1 0 0 0 0 0 0 0 0 1 0 1 0 1 0 0 1 0 1 0 0 1 1 0 0 1 0 1 1 0 0 0 0 0 1
[449] 1 1 0 0 0 0 0 0 0 1 1 1 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 1 0 1 0 0 1 0 0 0 0 0 0 1 0 1 1 1 0 1 1 0 0
[505] 0 1 1 0 0 0 0 1 1 0
```

```
> cat("AUC", fill = T)
AUC
> aucs
[1] 0.1520229 0.8479771
```



Thank you so much for your effort and help in this semester!

Thank you professor!